

CSC 656-01, S23

Extra Credit Coding Project

Due: Thu 13 Apr 2023 23:59 PDT – no late submissions

Overview

Using your implementation from [Coding Project #2](#), modify the `sum_vector` and `sum_indirect` codes as described below, run them on KNL nodes on `cori.nersc.gov` to collect new runtime information, and then use this new information to estimate the cache miss rate for each code at two different problem sizes.

Part 3 – Analyzing Results

Question 1. Please prepare a table with the following data (there are no charts in this project):

	# loops (integer)	Total memory bytes read (integer, divide by 1G for readability)	Total runtime (secs)	Avg Latency (nanoseconds)	Miss_rate (A % in the range 0..99.99)
Sum_vector 1M	3	0.0309	0.0297	1.0252	6.3388
Sum_vector 256M	3	0.0619	0.0582	0.9353	5.3147
Sum_indirect 1M	3	0.1238	0.1187	0.9567	5.2385
Sum_indirect 256M	3	0.2477	0.2396	0.9691	5.1954

Loops (integer) is the number of times your code called the `sum()` method in the >30 second testing window

Total memory bytes read: the total number of bytes your program loads from memory in the >30 second testing window

Total runtime of your program. It had better be > 30 secs. :)

Your average memory latency in *nanoseconds* computed as:

$\text{Total elapsed runtime in seconds} * 1,000,000,000 / \text{total bytes read}$

Note: your average memory latency is your Average Memory Access Time (AMAT) in nanoseconds

Miss_rate: computing this number is the point of this entire exercise :)

From L18 slide 48:

$$(AMAT - hit_time) / miss_penalty = miss_rate$$

Recall that

- Hit_time is a constant (see L18 slide 55)
- Miss_penalty is a constant (see L18 slide 55)
- Both hit_time and miss_penalty are time measurements in nanoseconds

Question 2. Consider the data in your table, and answer the following questions (3 sentences max for each question):

- a. Which code configuration shows the highest miss rate, and why?

The cache miss rate is affected by the memory access pattern of a program. If a program accesses memory in a contiguous pattern, the cache can store all the accessed data in its blocks, leading to a lower miss rate. However, if a program accesses memory in a non-contiguous pattern, the cache may have to evict data to make space for new data, leading to a higher miss rate.

- b. Which code configuration shows the lowest miss rate, and why?

The code configuration with the lowest miss rate is sum_direct. This is because it accesses contiguous memory locations in a linear fashion, which is more cache-friendly than accessing data in a scattered fashion. As a result, the cache can fetch and store data more efficiently, resulting in a lower miss rate.

For full credit, your response needs to include reasoning about the relationship between memory access patterns and its impact on the cache miss rate.